

# Delay and Rate-Optimal Control in a Multi-Class Priority Queue


Chih-ping Li

Massachusetts Institute of Technology

joint work with Michael Neely (USC)

INFOCOM 2012

# Revisit a classical queue control problem




An M/G/I queue serves  $N$  traffic classes

Pick a class, serve a job (nonpreemptively) one at a time

The service rate may be controllable (with costs)

Goal: Solve convex delay optimization problems

# Q1: Providing delay guarantees



## Feasibility


$$\overline{W}_1 \leq 2$$

$$\overline{W}_2 \leq 5$$

.....

$$\overline{W}_N \leq 3$$

## Q2: Providing delay fairness



**Fairly** treat two types of patients





Convex delay penalty

$$\text{Minimize: } \sum_{n=1}^N f_n(\overline{W}_n)$$

$$\text{subject to: } \overline{W}_n \leq d_n, \quad \forall n$$

## Q3: Minimizing service cost



(update service rates across busy periods)

### Rate control

Minimize:  $P_{av}$


subject to:  $\overline{W}_1 \leq 2$

.....

$\overline{W}_N \leq 3$

# Optimization approach

(i) Compute performance region




(ii) Solve optimization problem

$$\begin{aligned} &\text{Minimize: } \sum_{n=1}^N c_n \overline{W}_n \\ &\text{subject to: } (\overline{W}_n)_{n=1}^N \in \Lambda \end{aligned}$$

Gelenbe, Mitrani, Federgruen, Groenevelt, Shanthikumar, Tsoucas, Bertsimas, Dacre, Glazebrook, Garbe, Nino-mora, Yao, Coffman, .....

# Linear optimization is elegantly solved




$$\begin{aligned} \text{Minimize: } & \sum_{n=1}^N c_n \lambda_n \overline{W}_n \\ \text{subject to: } & (\overline{W}_n)_{n=1}^N \in \Lambda \end{aligned}$$

Delay region (with fixed rate)  
is the base of a polymatroid  
[Federgruen-Groenevelt '88]

**The  $c\mu$  rule:**  
Assign priorities to job classes in  
the decreasing order of  $c_n \mu_n$

## More good news



Linear program


$$\text{Minimize: } \sum_{n=1}^N c_n \overline{W}_n$$

$$\text{subject to: } (\overline{W}_n)_{n=1}^N \in \Lambda$$

Every vertex = a priority policy




However, in general cases....



$$\begin{aligned} \text{Minimize: } & \sum_{n=1}^N f_n(\overline{W}_n) \\ \text{subject to: } & (\overline{W}_n)_{n=1}^N \in \Lambda \\ & \overline{W}_n \leq d_n, \forall n \end{aligned}$$

Offline solution doesn't reveal the optimal scheduling policy

# Complex performance region



Could be difficult  
to describe

# A new methodology to solve these problems

## Feasibility

$$\overline{W}_1 \leq 2$$

$$\overline{W}_2 \leq 5$$

.....

$$\overline{W}_N \leq 3$$

## Convex delay penalty

$$\text{Minimize: } \sum_{n=1}^N f_n(\overline{W}_n)$$

$$\text{subject to: } \overline{W}_n \leq d_n, \forall n$$

## Rate control

$$\text{Minimize: } P_{av}$$


$$\text{subject to: } \overline{W}_1 \leq 2$$

.....



$$\overline{W}_N \leq 3$$

Optimal policies  
are as simple as  
**the  $c\mu$  rule**

## A new methodology




Performance region =  
a set of inequalities




Performance region =  
all **convex combinations**  
of vertices

## A new methodology





strict priority policy



strict priority policy  
with a fixed rate

## A new methodology



The optimal solution  
= an optimal time-sharing of vertex-achieving policies

# Providing delay guarantees

Minimize:  $I$

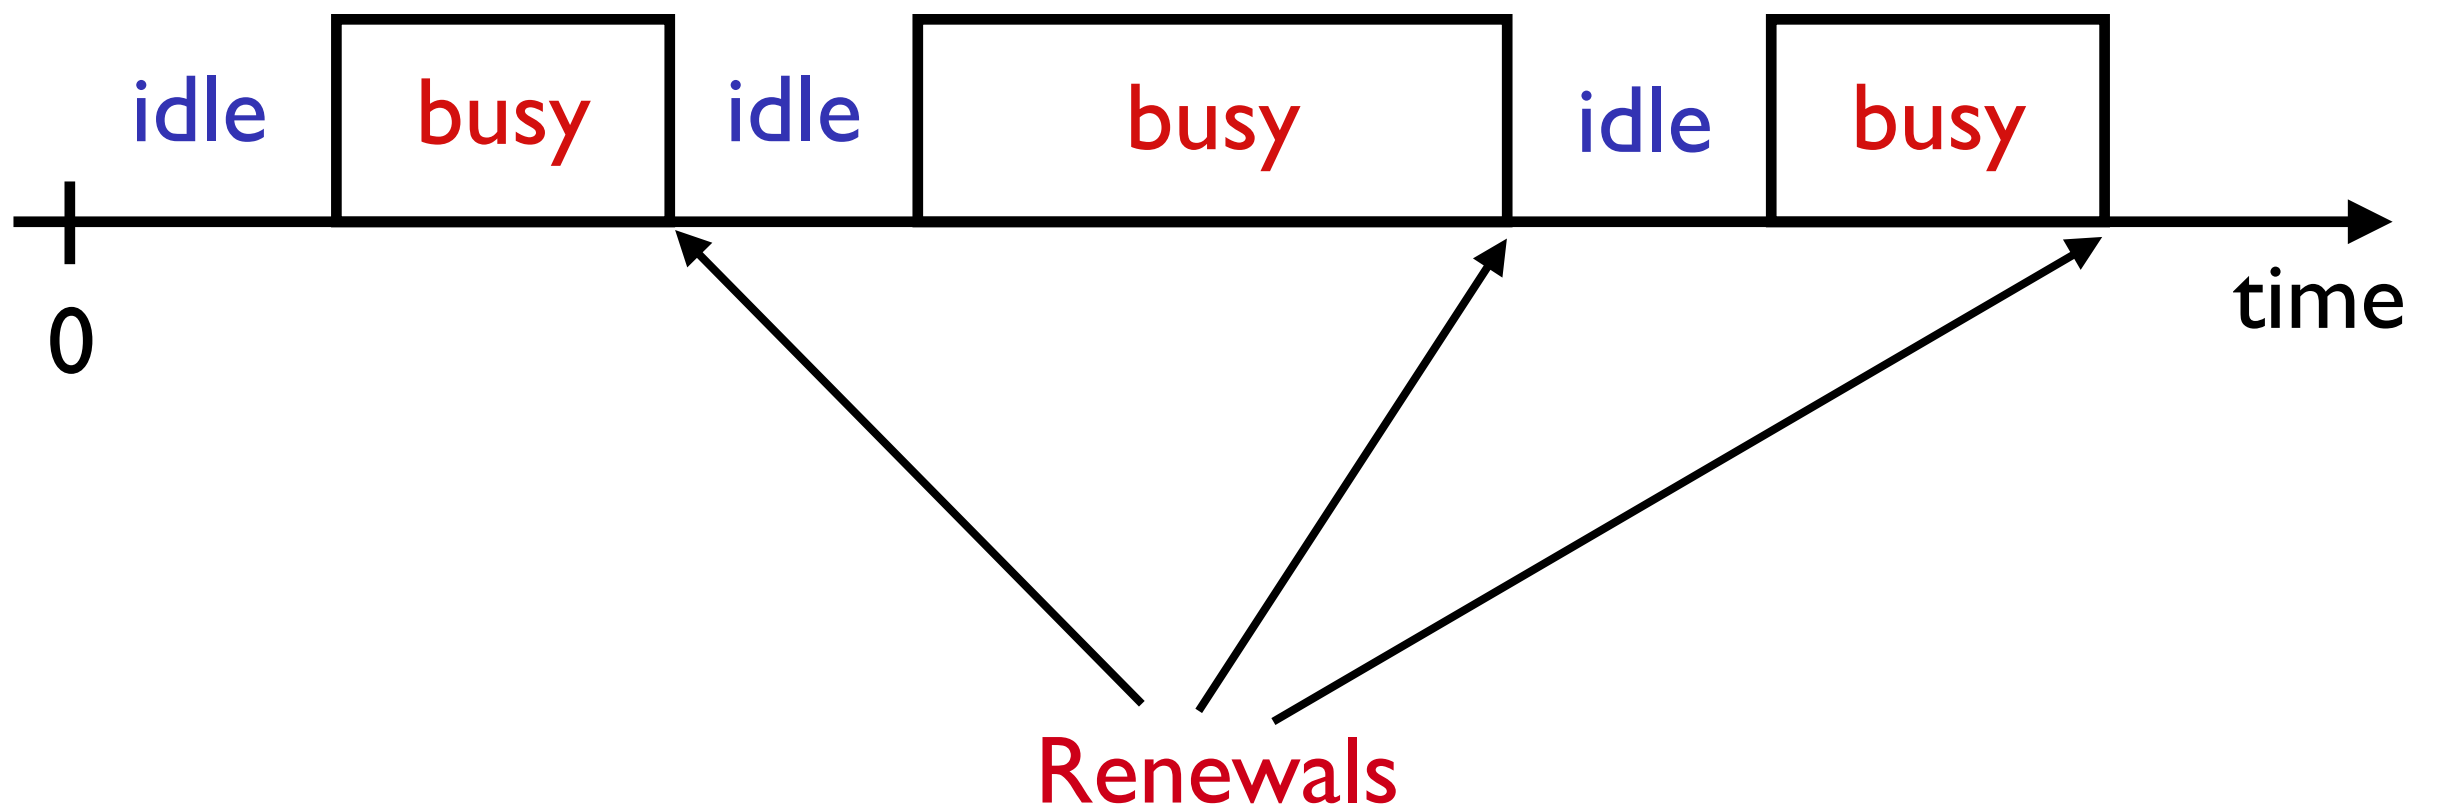
subject to:  $\overline{W}_n \leq d_n, n \in \{1, \dots, N\}$

$$(\overline{W}_n)_{n=1}^N \in \Lambda$$



# Providing delay guarantees



One priority policy **per busy (renewal) period**





# Providing delay guarantees

1. initial priority ordering



3. compute the delay debt owed to each class and re-prioritize


$$Z_1 \leftarrow \max [Z_1 + (8 - d_1) + (7 - d_1), 0]$$

$$Z_2 \leftarrow \max [Z_2 + (5 - d_2), 0]$$

$$Z_3 \leftarrow \max [Z_3 + (4 - d_3) + (2 - d_3), 0]$$

# Providing delay guarantees

Re-prioritize job classes in the decreasing order of  $Z_n$



$$Z_n \leftarrow \max \left[ Z_n + \sum_{i \in \text{class } n} (w_n^{(i)} - d_n), 0 \right]$$


amount of delay exceeding the desired bound

No statistical knowledge, low complexity

# Providing delay guarantees

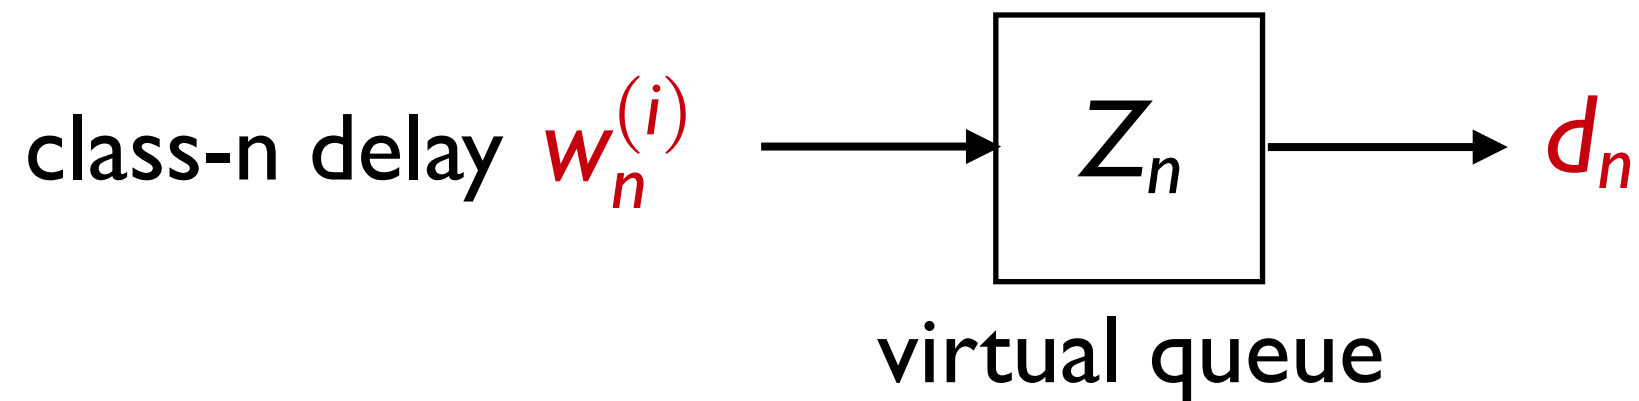
Theorem:

For every feasible  $\{d_1, d_2, \dots, d_n\}$ , we have  $\overline{W}_n \leq d_n$  for all  $n$ .



## Intuition

$$Z_n \leftarrow \max \left[ Z_n + \sum_{i \in \text{class } n} (w_n^{(i)} - d_n), 0 \right]$$



If queue  $Z_n$  is stable, then  $\overline{W_n} \leq d_n$

The dynamic priority policy stabilizes all virtual queues

(a max-weight policy in this context)


# Providing delay fairness

$$\text{Minimize: } \sum_{n=1}^N f_n(\overline{W}_n)$$

$$\text{subject to: } \overline{W}_n \leq d_n, \forall n$$



# Providing delay fairness



In every busy period, prioritize job classes by sorting  $\frac{Z_n + Y_n}{S_n}$

## Two virtual queues:

$$Z_n \leftarrow \max \left[ Z_n + \sum_{i \in \text{class } n} (w_n^{(i)} - d_n), 0 \right]$$

$$Y_n \leftarrow \max \left[ Y_n + \sum_{i \in \text{class } n} (w_n^{(i)} - \gamma_n), 0 \right]$$


$\gamma_n$  is the solution to:

Minimize:  $V f_n(\gamma_n) - Y_n \lambda_n \gamma_n$   
subject to:  $0 \leq \gamma_n \leq d_n$

## Intuition

$$Z_n \leftarrow \max \left[ Z_n + \sum_{i \in \text{class } n} (w_n^{(i)} - d_n), 0 \right]$$

$$Y_n \leftarrow \max \left[ Y_n + \sum_{i \in \text{class } n} (w_n^{(i)} - \gamma_n), 0 \right]$$



If queue  $Z_n$  is stable, then  $\overline{W}_n \leq d_n$

If queue  $Y_n$  is stable, then  $\overline{W}_n \leq \overline{\gamma}_n$

$$\sum_{n=1}^N f_n(\overline{W}_n) \leq \sum_{n=1}^N f_n(\overline{\gamma}_n)$$

Keep queues  $Z_n$  and  $Y_n$  stable, and minimize  $\sum_{n=1}^N f_n(\overline{\gamma}_n)$

# Providing delay fairness

Theorem:

For every feasible  $\{d_1, d_2, \dots, d_n\}$ , we have  $\overline{W}_n \leq d_n$  for all  $n$ .

$$\text{optimal} \leq \sum_{n=1}^N f_n(\overline{W}_n) \leq \text{optimal} + \frac{B}{V}$$

$V$  : a positive control parameter



# Simulation

2-class M/M/1 queue:


$$\text{Minimize: } \frac{1}{2} (\overline{W}_1)^2 + 2 (\overline{W}_2)^2$$

$$\text{subject to: } \overline{W}_1 + \overline{W}_2 = 2.4$$

$$0.4 \leq \overline{W}_1 \leq 1.95, \quad 0.4 \leq \overline{W}_2 \leq 1$$


$V$	$\overline{W}_1$	$\overline{W}_2$	Mean delay penalty
100	1.6607 (0.0055)	0.7424 (0.0052)	2.4814 (0.0239)
1000	1.7977 (0.0057)	0.5984 (0.0043)	2.3321 (0.0199)
2000	1.8339 (0.0056)	0.5639 (0.0053)	2.3176 (0.0217)
5000	1.8679 (0.0073)	0.5276 (0.0050)	2.3014 (0.0222)
Optimal value:	1.92	0.48	2.304

# Min-cost dynamic rate control



$$\begin{aligned} \text{Minimize: } & P_{av} \\ \text{subject to: } & \overline{W}_n \leq d_n, \quad \forall n \\ & P(t) \in [P_{\min}, P_{\max}] \end{aligned}$$

# Min-cost dynamic rate control



In every busy period,

Assign priorities by sorting  $\frac{Z_n}{S_n}$

$$Z_n \leftarrow \max \left[ Z_n + \sum_{i \in \text{class } n} (w_n^{(i)} - d_n), 0 \right]$$

Spend cost  $P$  minimizing

$$c V \frac{P}{\mu(P)} + \sum_{n=1}^N Z_n \lambda_n \overline{W}_n(P)$$

# Min-cost dynamic rate control

Theorem:

For every feasible  $\{d_1, d_2, \dots, d_n\}$ , we have  $\overline{W}_n \leq d_n$  for all  $n$ .

$$\text{optimal} \leq P_{av} \leq \text{optimal} + \frac{B}{V}$$

$V$  : a positive control parameter

# Min-cost dynamic rate control


2-class M/G/1 queue

$$\mu(P) = \sqrt{P}, P(t) \in \{16, 25\}$$

Minimize:  $\bar{P}$

subject to:  $\bar{W}_1 \leq 0.4$

$\bar{W}_2 \leq 0.325$



$V$	$\bar{W}_1^{\text{DynRate}}$	$\bar{W}_2^{\text{DynRate}}$	Average cost
1	0.3562 (0.00078)	0.3029 (0.00032)	13.8018 (0.01806)
10	0.3984 (0.00022)	0.3247 (0.00005)	13.5101 (0.02626)
100	0.4003 (0.00013)	0.3252 (0.00010)	13.5044 (0.02197)
Optimal value:	0.4	0.325	13.5

# Dynamic index policies

Delay  
guarantees

Delay  
Fairness

Rate  
control

Index

$$Z_n$$

$$(Z_n + Y_n) / \bar{S}_n$$

$$Z_n / \bar{S}_n$$

Statistical  
knowledge

none

$$\lambda_n, \bar{S}_n$$

$$\lambda_n, \bar{S}_n, \bar{W}_n(P)$$

$$Z_n \leftarrow \max \left[ Z_n + \sum_{i \in \text{class } n} (w_n^{(i)} - d_n), 0 \right]$$

$$Y_n \leftarrow \max \left[ Y_n + \sum_{i \in \text{class } n} (w_n^{(i)} - \gamma_n), 0 \right]$$

# Summary

Convex programs over a polymatroid can be solved by online adaptive policies as simple as the  $c\mu$  rule

These policies require limited or no statistics by learning from the past

Caveat: large delay variance when queue is heavily loaded (controls are updated over busy periods)

Applications to other queueing systems with polymatroid-type performance regions?

Job-level scheduling using imperfect virtual-queue length information to reduce delay variance?