

Receiver-Based Flow Control for Networks in Overload

Chih-ping Li


Massachusetts Institute of Technology

joint work with Eytan Modiano (MIT)

IEEE INFOCOM April 18, 2013

Traditional Flow Control


- Sources adjust transmission rates in response to network congestion
 - TCP flow control
 - Network utility maximization problems



$$\begin{aligned} &\text{maximize: } \sum U_n(x_n) \\ &\text{subject to: } (x_n) \in \Lambda \end{aligned}$$

Not All Sources Perform Flow Control


- Greedy users
- Malicious users launch attacks to crash prominent websites
- UDP-based flows



If flow 2 overloads the network, then flows 1 & 3 are adversely affected

Not All Sources Perform Flow Control

- Greedy users
- Malicious users launch attacks to crash prominent websites
- UDP-based flows




If flow 2 overloads the network, then flows 1 & 3 are adversely affected

Task 1:

In-network packet dropping by all intermediate network nodes to optimize per-flow throughput without source cooperation


When Source-Based Flow Control is Ineffective

- Flash crowd - lots of small flows yield congestion close to a web server (e.g., news websites experience high delay after the 9/11 attack)



When Source-Based Flow Control is Ineffective

- Flash crowd - lots of small flows yield congestion close to a web server (e.g., news websites experience high delay after the 9/11 attack)




Task 2:

Let the receiver control **the aggregate rate** of a class of flows without source cooperation

When Source-Based Flow Control is Ineffective

- Flash crowd - lots of small flows yield congestion close to a web server (e.g., news websites experience high delay after the 9/11 attack)




Task 2:

Let the receiver control **the aggregate rate** of a class of flows without source cooperation

Provide fairness (differentiated services) **over classes of flows**, instead of individual flows

Task I: Maximizing weighted sum throughput
by in-network packet dropping

Maximum Weighted Sum Throughput




Unit capacity links
Arrival rates (2, 2, 2)

r_i - long-term throughput of flow i

Maximize: $3 r_1 + 2 r_2 + 1 r_3 \longrightarrow$ Optimal throughput (1, 0, 1)

B serves flow 1, A serves flow 3

Maximum Weighted Sum Throughput



Unit capacity links
Arrival rates (2, 2, 2)

r_i - long-term throughput of flow i


Maximize: $3 r_1 + 2 r_2 + 1 r_3 \longrightarrow$ Optimal throughput (1, 0, 1)

B serves flow 1, A serves flow 3


Maximize: $3 r_1 + 5 r_2 + 1 r_3 \longrightarrow$ Optimal throughput (0, 1, 0)

A & B serve flow 2

Threshold-Based Packet Dropping




Each node has a network-layer queue $Q(t)$ and a drop queue $D(t)$ for each flow




The use of $D(t)$: directly control average packet dropping rate


Threshold-Based Packet Dropping



Dropping policy for flow 2 at node A:




Threshold-Based Packet Dropping




Maximize: $\theta_1 r_1 + \theta_2 r_2 + \theta_3 r_3$

Dropping policy for **flow 2** at node A:




d_{\max} : a finite constant guaranteeing the network can be stabilized by packet dropping

Threshold-Based Packet Dropping




Dropping policy for **flow 2** at node A:




d_{\max} : a finite constant guaranteeing the network can be stabilized by packet dropping

Overload Resilient Algorithm



Maximize: $\theta_1 r_1 + \theta_2 r_2 + \theta_3 r_3$



+

Back-pressure routing
over queues $Q_c(t)$

Theorem:

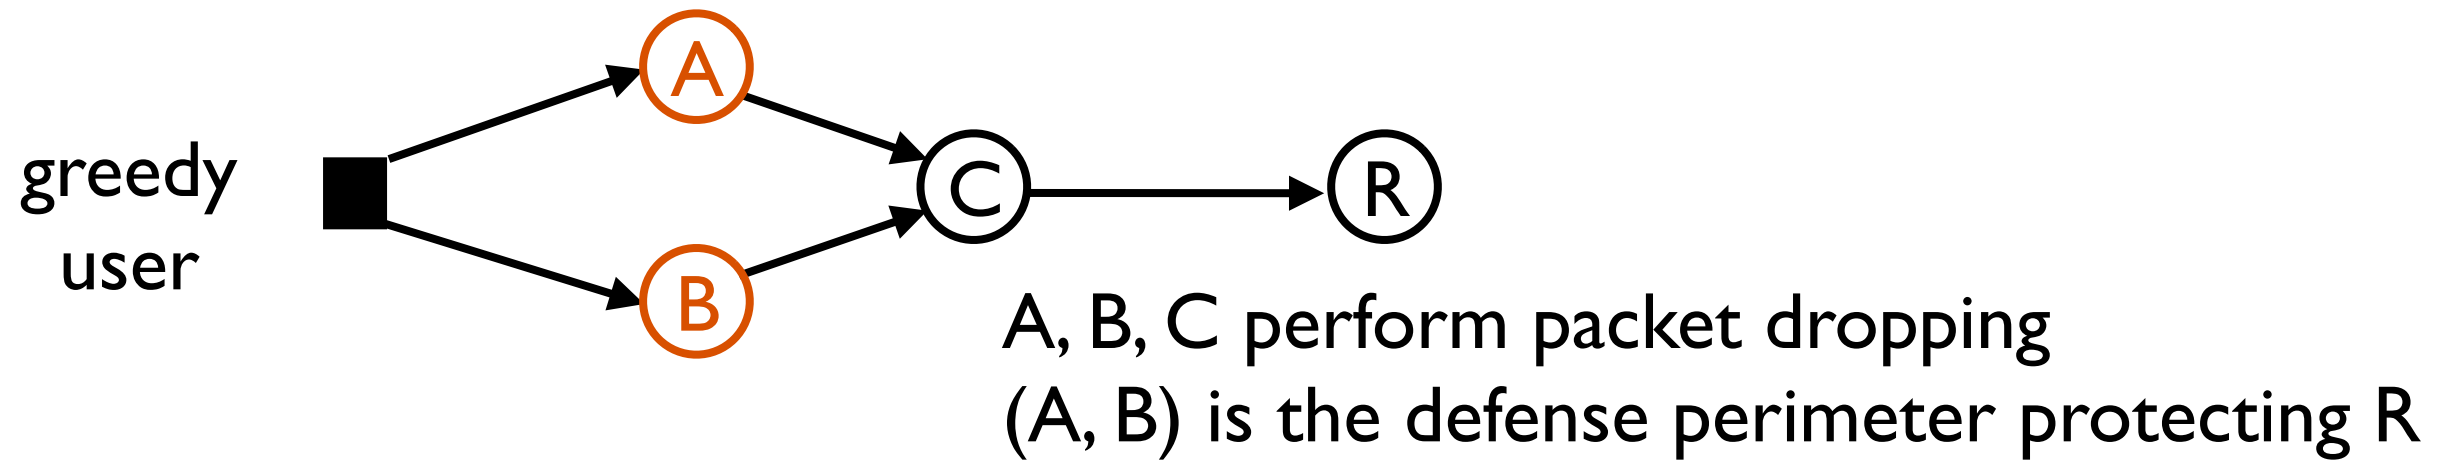
$Q_c(t) \leq V\theta_c + 2d_{\max}$ for all t
(using finite-size buffer $\propto \theta_c$)

$\sum \theta_c r_c \geq \text{optimal} - O(1/V)$
(performance gap diminishes
as buffer size increases)

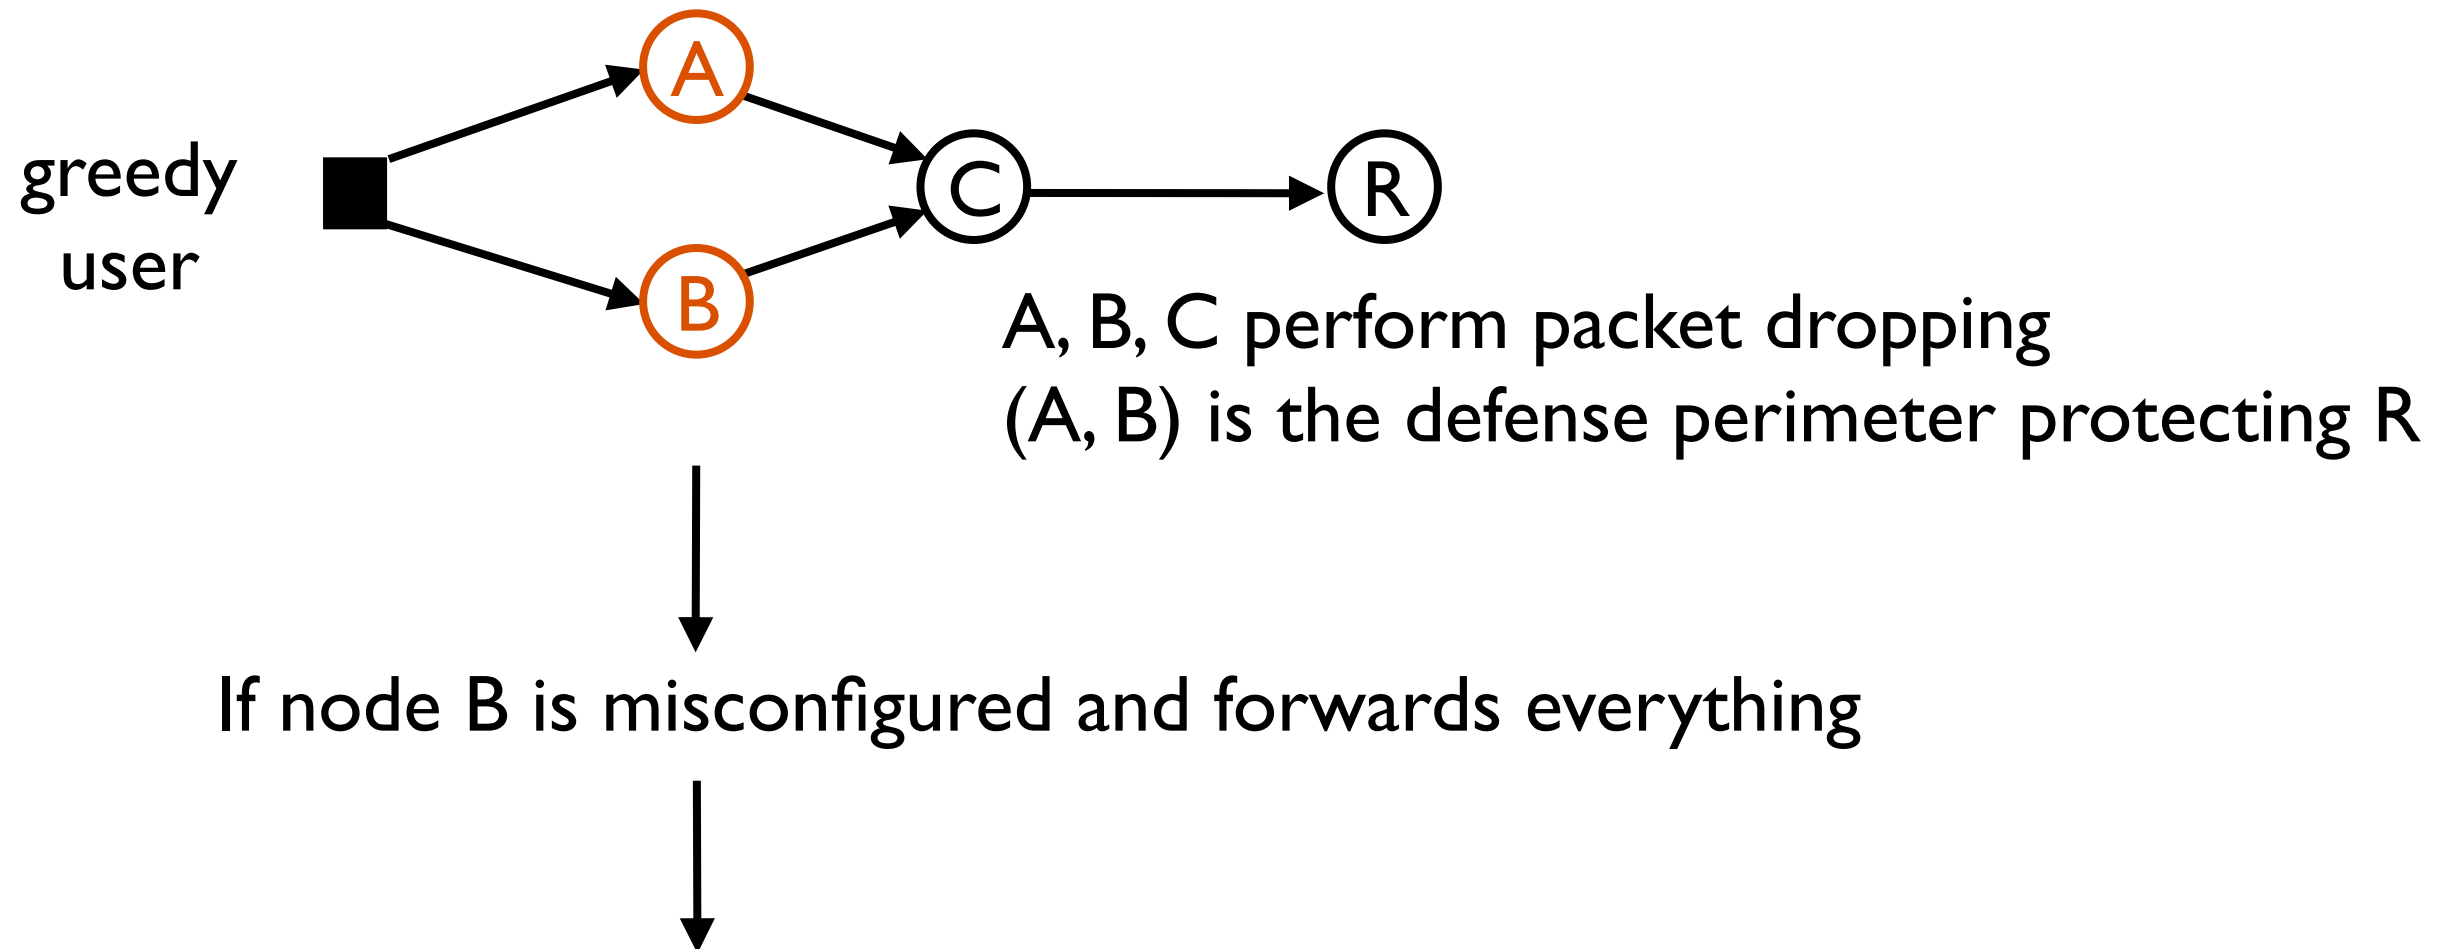
Independent of arrival rates

Distributed algorithm


In-Network Packet Dropping is Robust



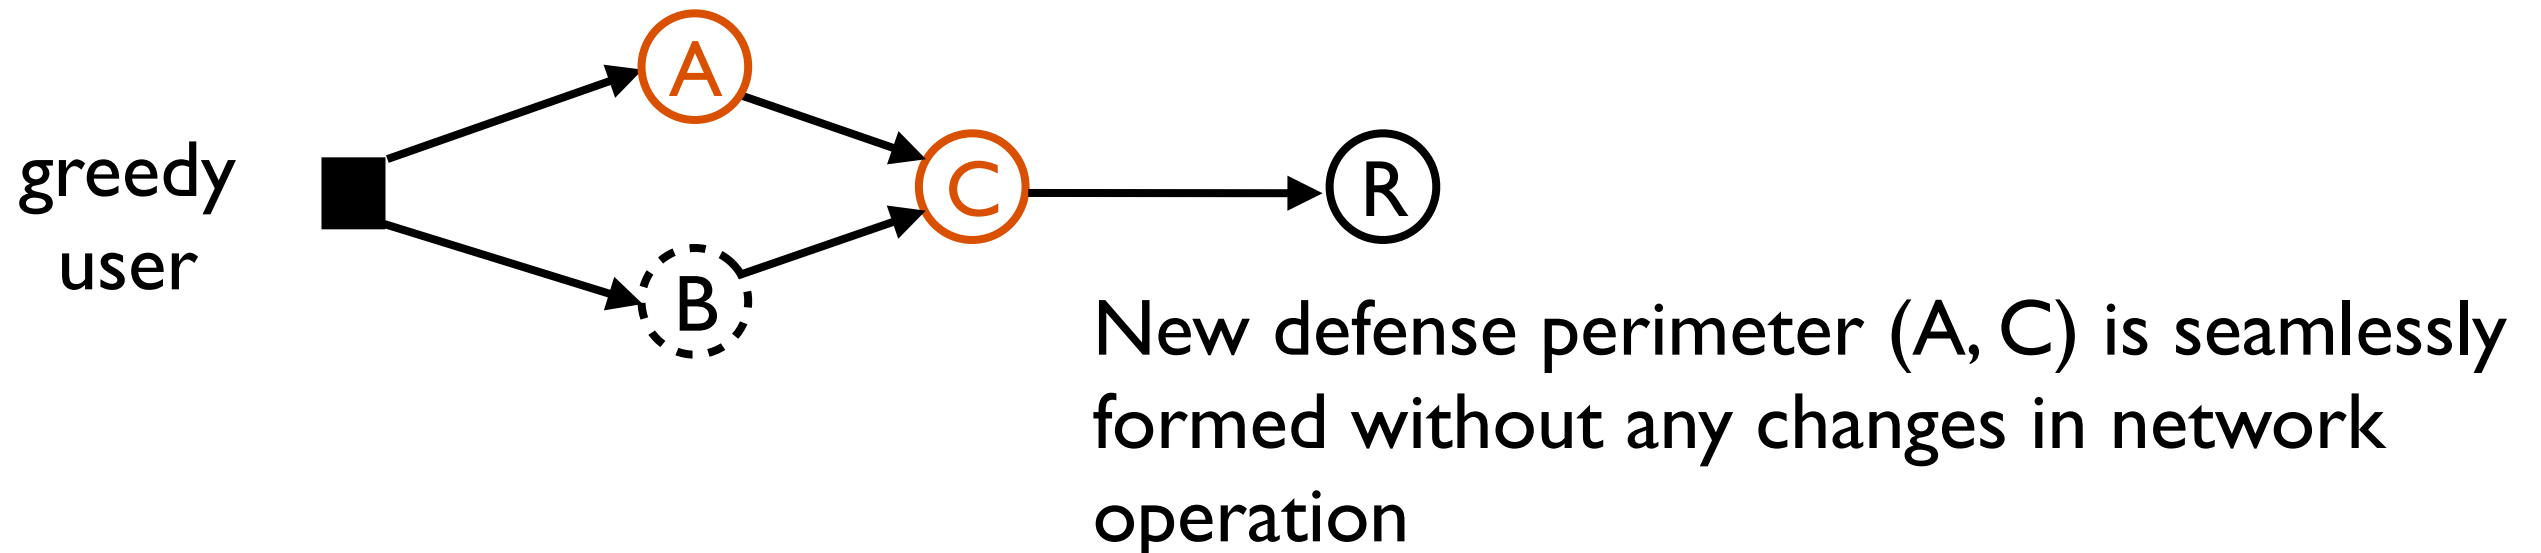
In-Network Packet Dropping is Robust



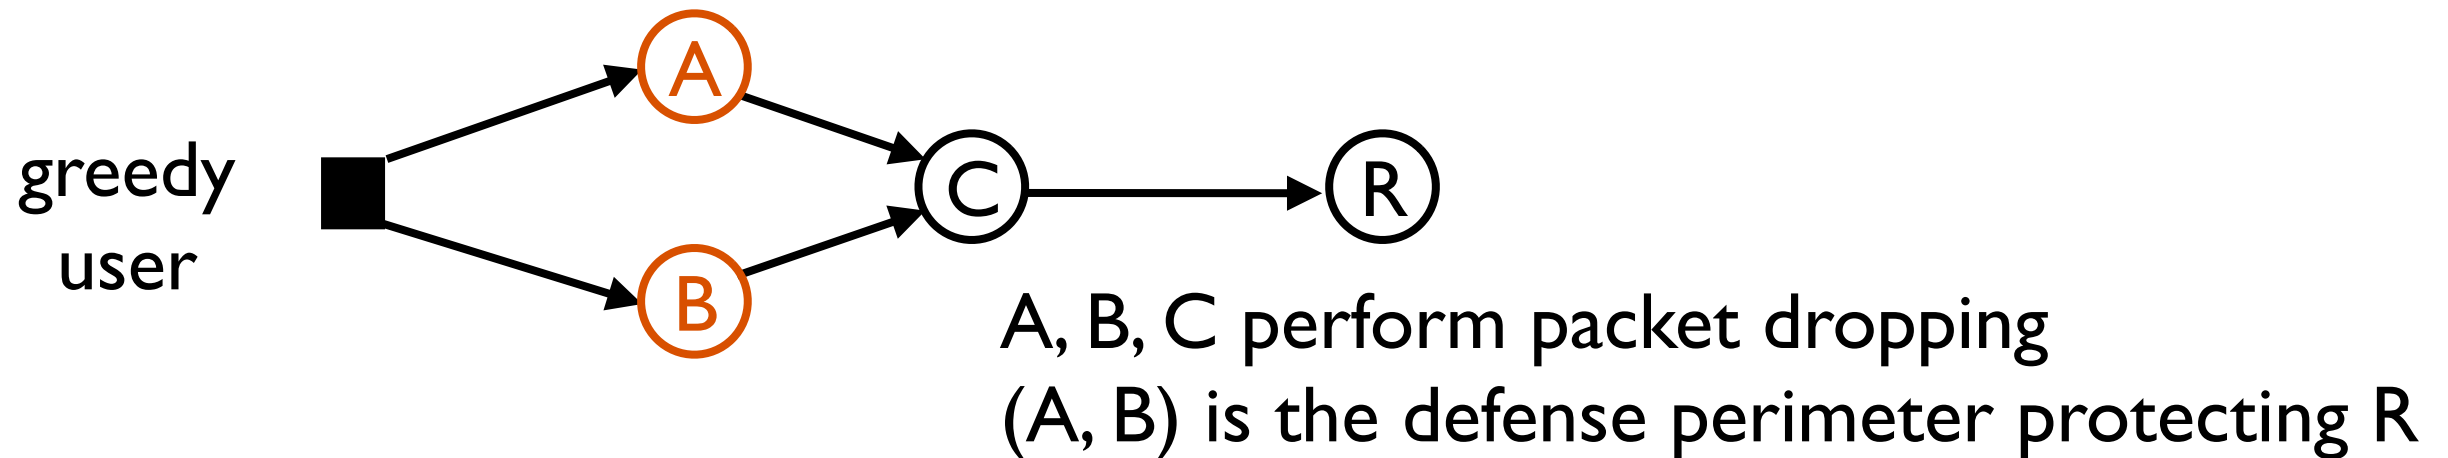
In-Network Packet Dropping is Robust



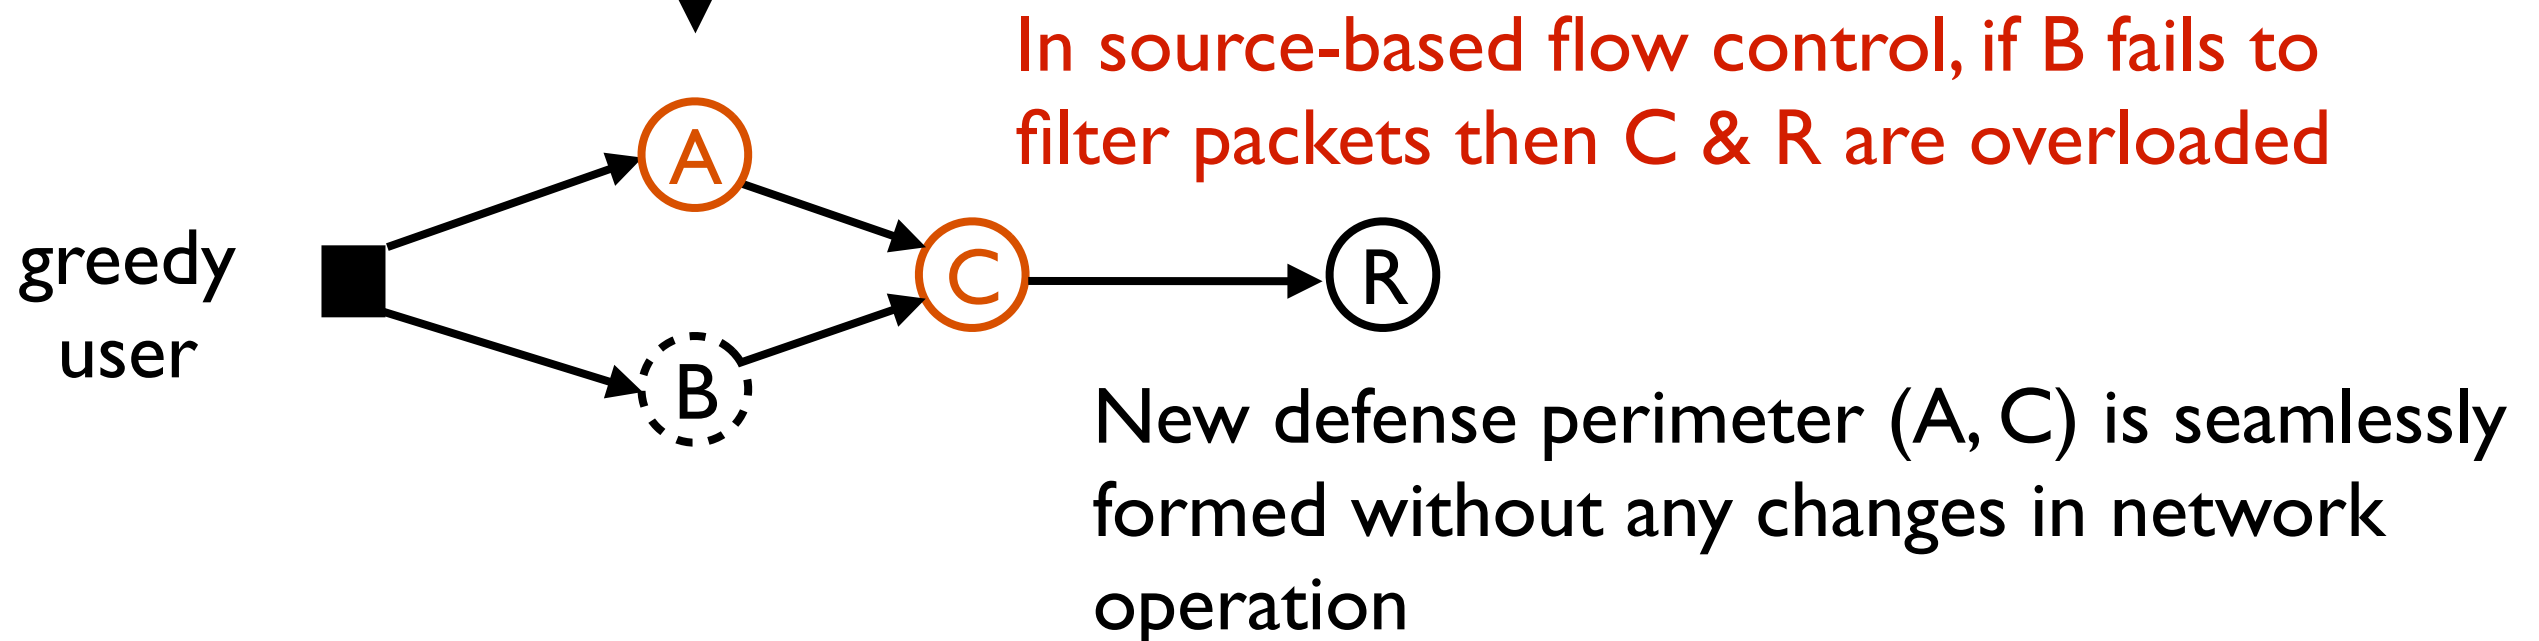
If node B is misconfigured and forwards everything




In-Network Packet Dropping is Robust



If node B is misconfigured and forwards everything



Simulation



Unit capacity links
Arrival rates (2, 2, 2)

(a) Maximizing $3r_1 + 2r_2 + r_3$

V	r_1	r_2	r_3
10	.787	.168	.099
20	.867	.133	.410
50	.992	.008	.967
100	.999	0	.999
opt	1	0	1

(b) Maximizing $3r_1 + 5r_2 + r_3$


V	r_1	r_2	r_3
10	.185	.815	.083
20	.107	.893	.095
50	.031	.969	.031
100	.002	.998	.001
opt	0	1	0

As buffer size
increases



Buffer size is $V\theta_c + 2d_{\max}$


Simulation



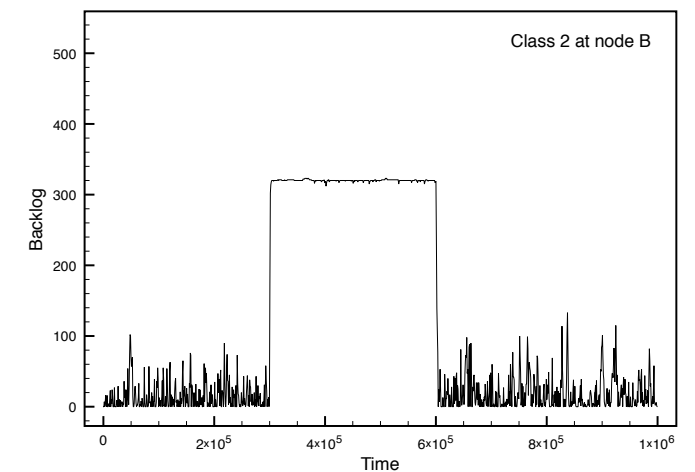
Time-varying arrival rates
 $(0.8, 0.1, 0.8) \rightarrow (0.8, 2, 0.8) \rightarrow (0.8, 0.1, 0.8)$
stable **overload** **stable**

Maximize $3 r_1 + 5 r_2 + r_3$

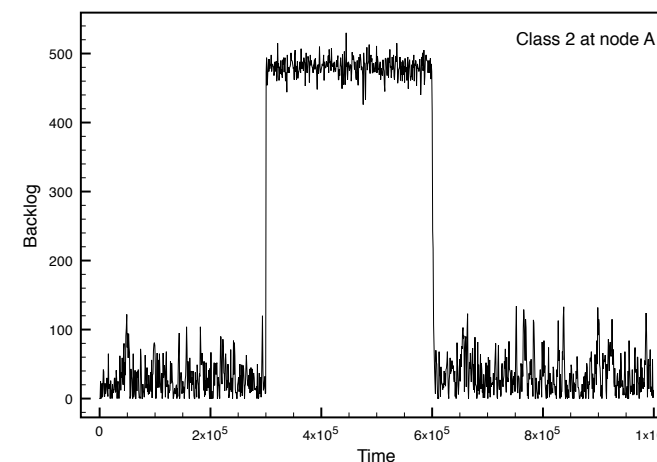
time interval	throughput in this interval	optimal value
$[0, 3 \cdot 10^5)$	$(.807, .104, .8)$	$(.8, .1, .8)$
$[3 \cdot 10^5, 6 \cdot 10^5)$	$(.002, .998, 0)$	$(0, 1, 0)$
$[6 \cdot 10^5, 10^6)$	$(.793, .106, .796)$	$(.8, .1, .8)$




(a) Queue $Q_B^{(1)}(t)$



(b) Queue $Q_B^{(2)}(t)$




(c) Queue $Q_A^{(2)}(t)$



(d) Queue $Q_A^{(3)}(t)$

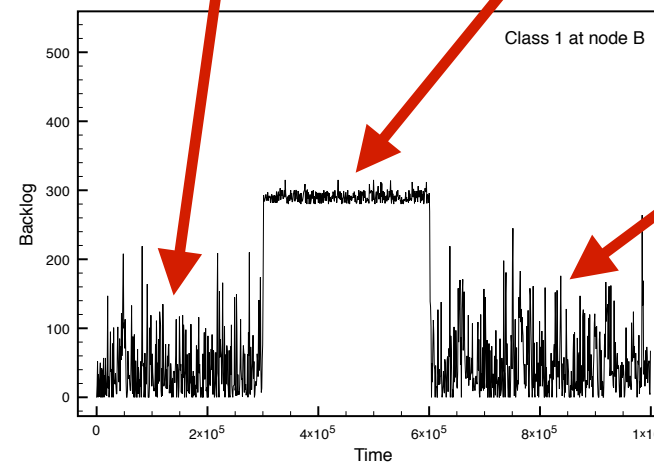
Simulation




Maximize $3 r_1 + 5 r_2 + r_3$

time interval	throughput in this interval	optimal value
$[0, 3 \cdot 10^5)$	$(.807, .104, .8)$	$(.8, .1, .8)$
$[3 \cdot 10^5, 6 \cdot 10^5)$	$(.002, .998, 0)$	$(0, 1, 0)$
$[6 \cdot 10^5, 10^6)$	$(.793, .106, .796)$	$(.8, .1, .8)$


Time-varying arrival rates
 $(0.8, 0.1, 0.8) \rightarrow (0.8, 2, 0.8) \rightarrow (0.8, 0.1, 0.8)$
 stable overload stable



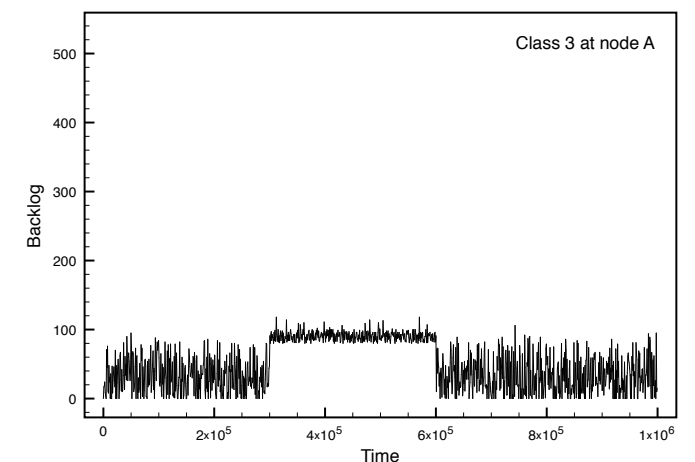
(a) Queue $Q_B^{(1)}(t)$



(b) Queue $Q_B^{(2)}(t)$




(c) Queue $Q_A^{(2)}(t)$



(d) Queue $Q_A^{(3)}(t)$

Task 2: Controlling the aggregate rate of flows by receiver-based flow control

Class-Based Utility Maximization




Assign a utility to a class of flows to control its aggregate throughput

$$\begin{aligned} &\text{maximize} && g_1(r_{1A} + r_{1C}) + g_2(r_{2B}) + g_3(r_{3D}) \\ &\text{subject to} && r_{ij} \leq \lambda_{ij}, (r_{1A}, r_{1C}, r_{2B}, r_{3D}) \text{ feasible} \end{aligned}$$

Generalizing the traditional approach of optimizing per-flow utilities

Class-Based Utility Maximization




$$\begin{aligned} &\text{maximize} && g_1(r_{1A} + r_{1C}) + g_2(r_{2B}) + g_3(r_{3D}) \\ &\text{subject to} && r_{ij} \leq \lambda_{ij}, \quad (r_{1A}, r_{1C}, r_{2B}, r_{3D}) \text{ feasible} \end{aligned}$$

Can we design a distributed control policy?

- non-separable objective function
- how to distributively drop packets, without knowing arrival rates, to optimize aggregate throughput of each class?


Receiver-End Flow Control




Control valve $Z_c(t)$ interacts with back-pressure routing to optimize class-c aggregate throughput

Utility-Optimal Overload Resilient Algorithm

- (I) Back-pressure routing, where each node maintains per-class queues.
For flow c over link (a, d_c) , use the differential backlog




$$W_l^{(c)}(t) = Q_a^{(c)}(t) - Q_{d_c}^{(c)}(t)$$




Utility-Optimal Overload Resilient Algorithm

- (I) Back-pressure routing, where each node maintains per-class queues.
For flow c over link (a, d_c) , use the differential backlog




$$W_l^{(c)}(t) = Q_a^{(c)}(t) - \underbrace{Q_{d_c}^{(c)}(t)}$$

=0 in normal
back-pressure




Utility-Optimal Overload Resilient Algorithm

- (I) Back-pressure routing, where each node maintains per-class queues.
For flow c over link (a, d_c) , use the differential backlog




$$w_l^{(c)}(t) = Q_a^{(c)}(t) - Q_{d_c}^{(c)}(t)$$




Utility-Optimal Overload Resilient Algorithm

- (I) Back-pressure routing, where each node maintains per-class queues.
For flow c over link (a, d_c) , use the differential backlog



$$w_l^{(c)}(t) = Q_a^{(c)}(t) - Q_{d_c}^{(c)}(t)$$




$$\text{if } Z_c(t) \geq Q \longrightarrow Q_{d_c}^{(c)}(t) = we^{w(Z_c(t)-Q)} > 0$$

(receiving too many)


Utility-Optimal Overload Resilient Algorithm

- (I) Back-pressure routing, where each node maintains per-class queues.
For flow c over link (a, d_c) , use the differential backlog



$$w_l^{(c)}(t) = Q_a^{(c)}(t) - Q_{d_c}^{(c)}(t)$$

push-back mechanism (slow down)




$$\text{if } Z_c(t) \geq Q \longrightarrow Q_{d_c}^{(c)}(t) = w e^{w(Z_c(t) - Q)} > 0$$


(receiving too many)

Utility-Optimal Overload Resilient Algorithm

- (I) Back-pressure routing, where each node maintains per-class queues.
For flow c over link (a, d_c) , use the differential backlog



$$W_l^{(c)}(t) = Q_a^{(c)}(t) - Q_{d_c}^{(c)}(t)$$




$$\text{if } Z_c(t) \geq Q \longrightarrow Q_{d_c}^{(c)}(t) = we^{w(Z_c(t)-Q)} > 0$$


(receiving too many)

Utility-Optimal Overload Resilient Algorithm

- (I) Back-pressure routing, where each node maintains per-class queues.
For flow c over link (a, d_c) , use the differential backlog



$$W_l^{(c)}(t) = Q_a^{(c)}(t) - Q_{d_c}^{(c)}(t)$$



$$\text{if } Z_c(t) \geq Q \longrightarrow Q_{d_c}^{(c)}(t) = we^{w(Z_c(t)-Q)} > 0$$


(receiving too many)

$$\text{if } Z_c(t) < Q \longrightarrow Q_{d_c}^{(c)}(t) = -we^{w(Q-Z_c(t))} < 0$$

(receiving too few)


Utility-Optimal Overload Resilient Algorithm

- (I) Back-pressure routing, where each node maintains per-class queues.
For flow c over link (a, d_c) , use the differential backlog



$$w_l^{(c)}(t) = Q_a^{(c)}(t) - Q_{d_c}^{(c)}(t)$$

“receiver asks for data” (speed up)



$$\text{if } Z_c(t) \geq Q \longrightarrow Q_{d_c}^{(c)}(t) = we^{w(Z_c(t)-Q)} > 0$$

(receiving too many)

$$\text{if } Z_c(t) < Q \longrightarrow Q_{d_c}^{(c)}(t) = -we^{w(Q-Z_c(t))} < 0$$

(receiving too few)


Utility-Optimal Overload Resilient Algorithm

(2) Choose $\nu_c(t)$ that solves

$$\text{maximize } V h_c(\nu_c(t)) + \nu_c(t) Q_{d_c}^{(c)}(t)$$

$$\text{subject to } 0 \leq \nu_c(t) \leq \nu_{\max}$$

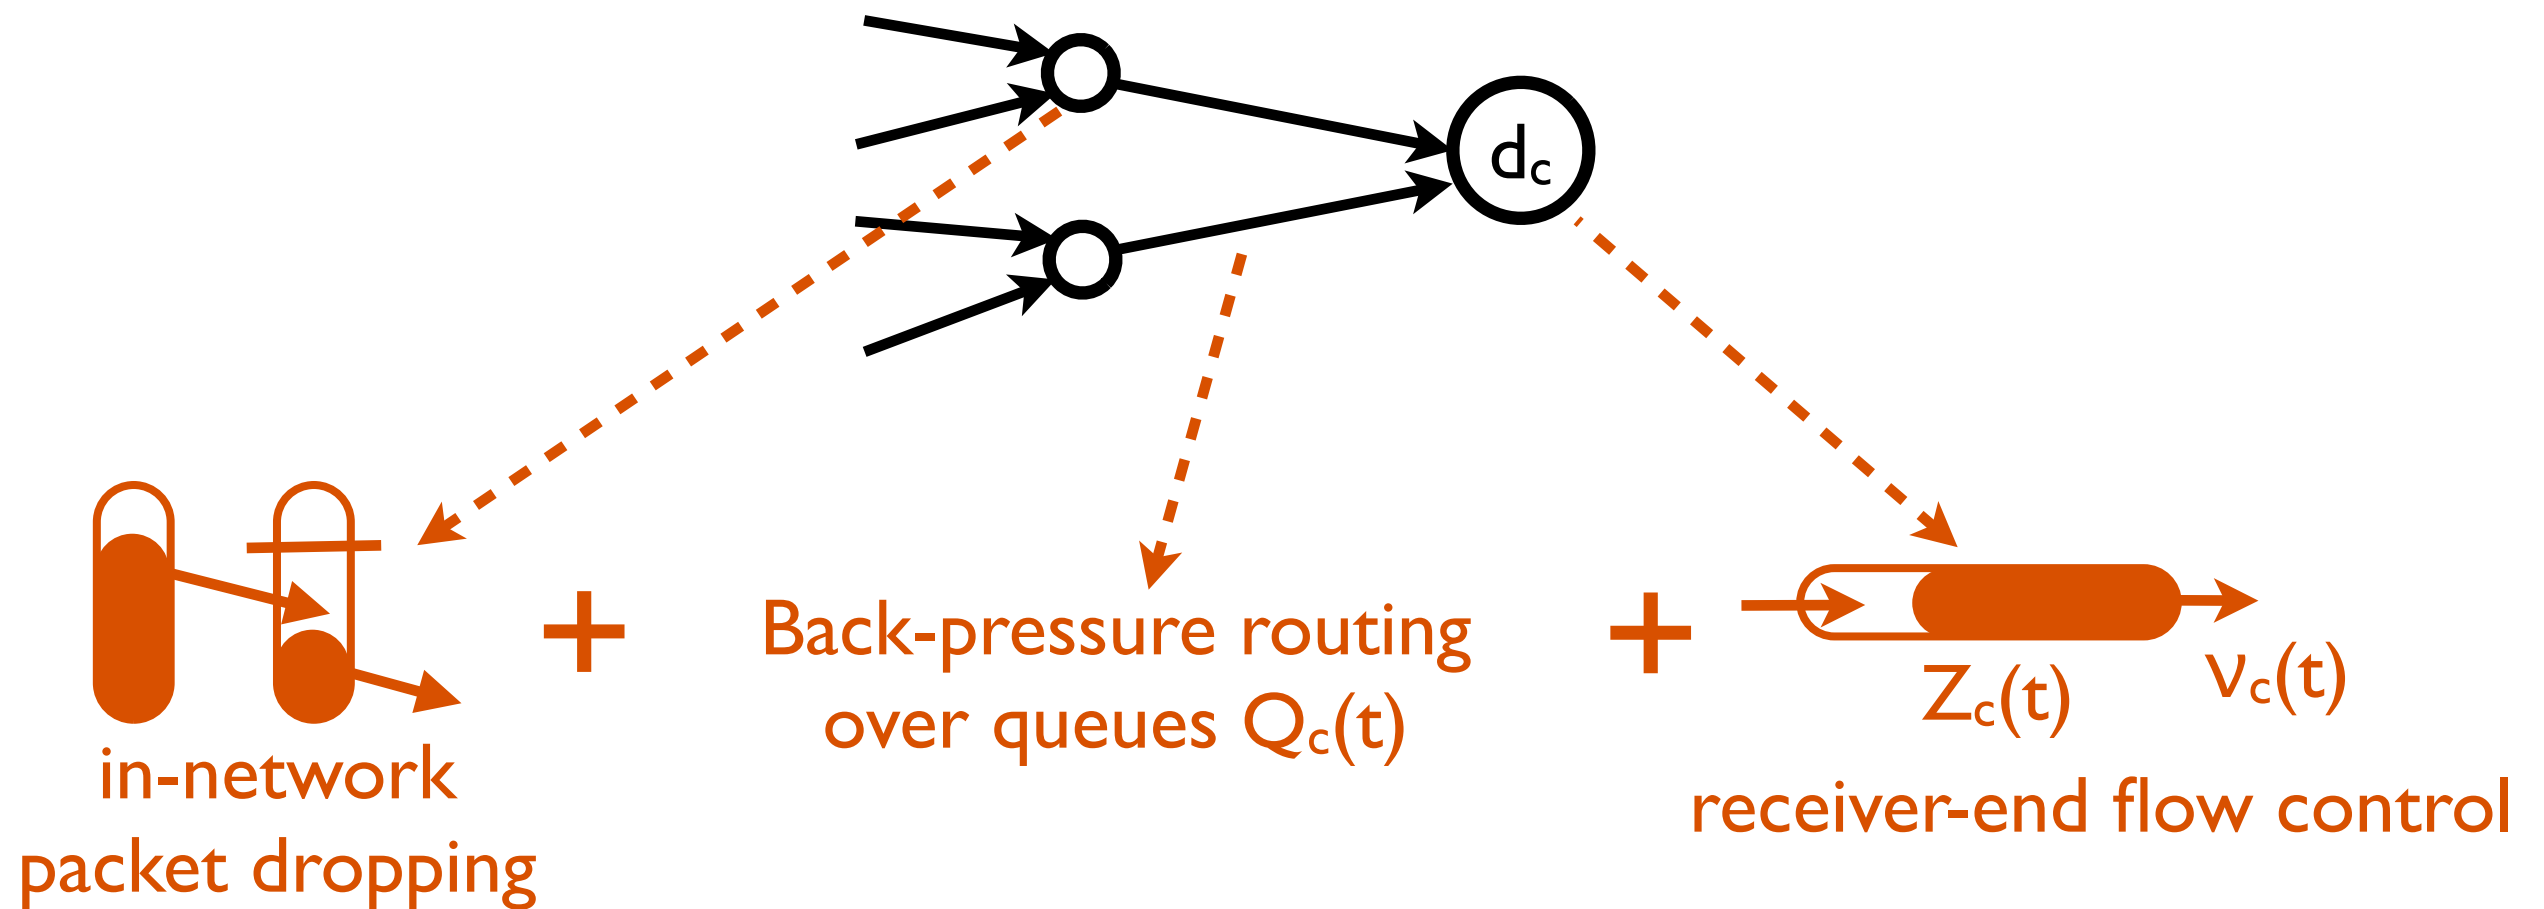
$$h_c(x) = g_c(x) - \theta_c x$$



(3) The same packet dropping policy for per-class queues at each node



Utility-Optimal Overload Resilient Algorithm



Theorem:

$$Q_c(t) \leq V\theta_c + 2d_{\max} \text{ for all } t \text{ (using finite buffer)}$$


$$\sum g_c(r_c) \geq \text{optimal utility} - O(1/V + \varepsilon)$$

(large buffer reduces performance gap)

Sketch of analysis

- Given a fixed arrival rate vector, characterizing the set of achievable throughput vectors in terms of queue overflow rates
 - optimal queue overflow rate = optimal packet dropping rate
- Decoupling a utility function into two components
 - new utility functions for receiver-end flow control
 - distributed packet dropping
 - optimizing original utility = max. new utility + min. packet dropping
- Transforming utility optimization problems using auxiliary variables
- Using exponential Lyapunov functions for drift analysis

Utility-Optimal Algorithm / Simulations



Unit capacity links

Arrival rates (2, 2, 2)

Objective: proportional fairness

Optimal throughput: (2/3, 1/3, 2/3)

Throughput


V	r_1	r_2	r_3	$\sum \log(r_c)$
10	.522	.478	.522	-2.038
20	.585	.415	.585	-1.952
50	.631	.369	.631	-1.918
100	.648	.352	.647	-1.912
optimal	.667	.333	.667	-1.910

Maximum backlog

V	$Q_B^{(1)}(t)$	$Q_B^{(2)}(t)$	$Q_A^{(2)}(t)$	$Q_A^{(3)}(t)$	$Q_{\max}^{(c)}$
10	140	97	137	137	142
20	237	187	240	236	242
50	539	441	538	540	542
100	1036	865	1039	1039	1042

Finite buffer size: $V\theta_c + 2d_{\max}$

Utility-Optimal Algorithm / Simulations



Unit capacity links
Arrival rates (2, 2, 2, 2)

Objective: max-min fairness **over classes**


Per-class throughput

V	$r_{1A} + r_{1C}$	r_{2B}	r_{3D}
10	.200	.100	.100
20	.364	.206	.205
30	.661	.650	.651
50	.667	.667	.667
optimal	.667	.667	.667


As buffer size
increases

Finite buffer size: $V\theta_c + 2d_{\max}$

Remarks

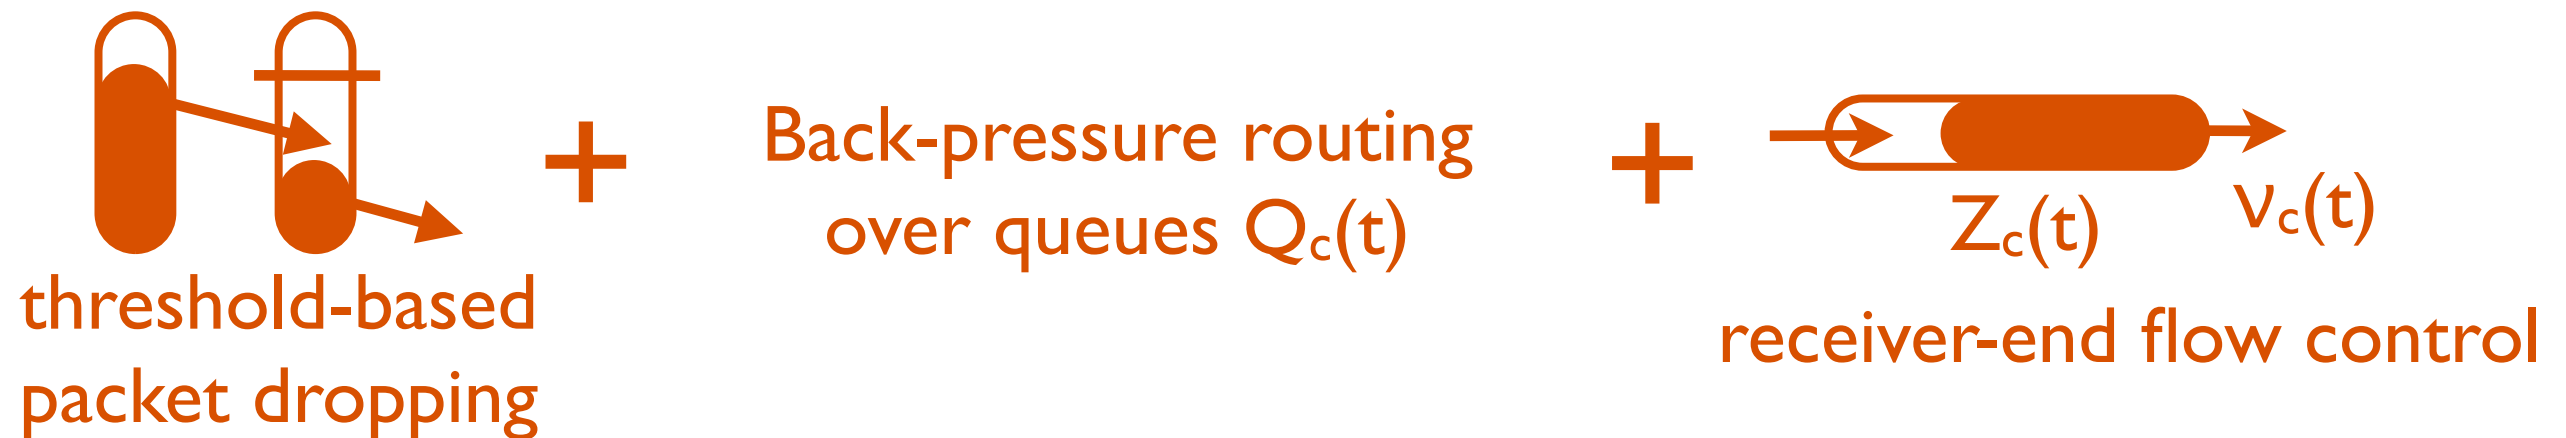


Remarks



Summary

- Key problem: optimizing network performance without source cooperation
- Can't rely on the sources to resolve network overload
 - greedy users, congestion-unaware traffic, flash crowd, DDoS attacks
- We propose a robust, distributed, utility-optimal algorithm:



- This algorithm controls the aggregate rate of flows, providing differentiated services to different classes of flows